

Introduction to Bayesian analysis 1, Take-home task  
Tuomas Porkamaa

1.

Likelihood:

$$Y = 12 \mid \theta \sim \text{Binom}(n = 20, \theta), p(Y = 12 \mid \theta) = \binom{20}{12} \theta^{12} (1 - \theta)^8$$

a.)

Prior:

$$\theta \sim \text{Unif}(0, 1)$$

Posterior:

$$p(\theta \mid Y = 12) \propto p(\theta) p(Y = 12 \mid \theta) \propto \theta^{12} (1 - \theta)^8$$

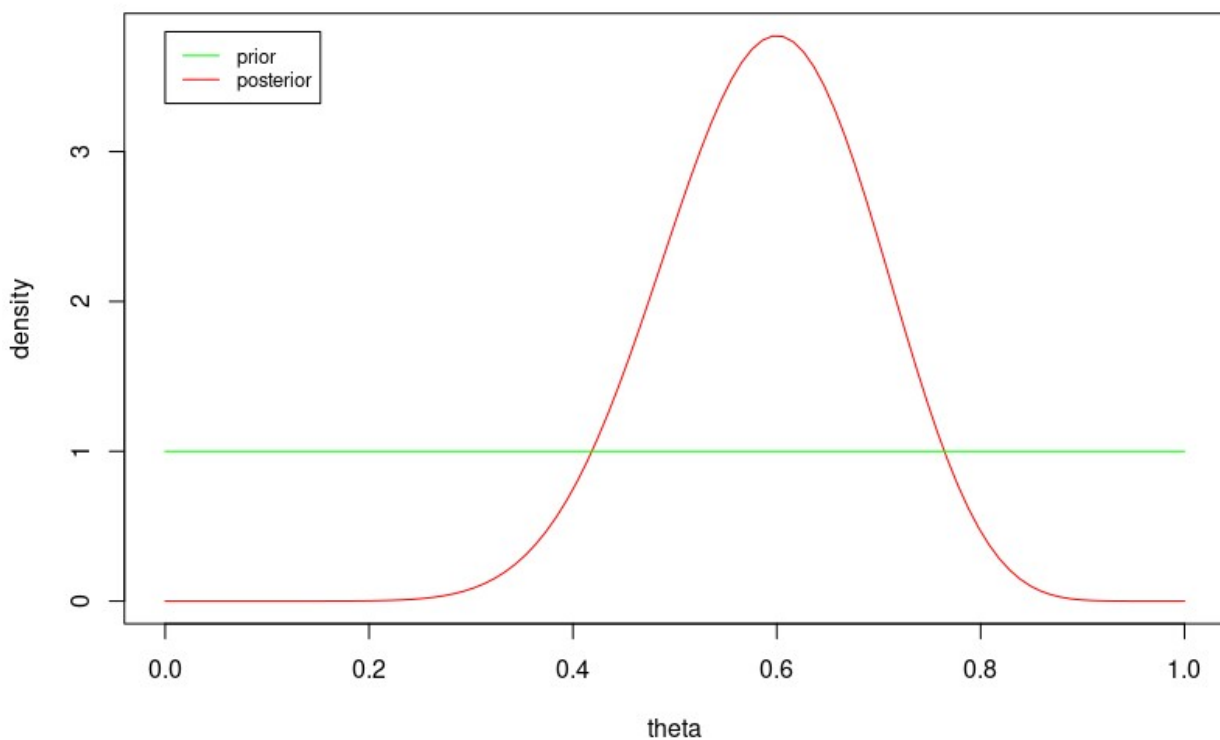
$$\therefore \theta \mid Y = 12 \sim \text{Beta}(13, 9)$$

$$\therefore E(\theta \mid Y = 12) = 0.59, \text{Var}(\theta \mid Y = 12) = 0.0105$$

**R-Codes and plots:**

```
theta = seq(0,1,length.out = 100)
thetaPostDist <- dbeta(theta, 13, 9)
thetaPriorDist <- dunif(theta, 0, 1)
```

```
plot(theta, thetaPostDist, type = "l", col="red", ylab="density")
lines(theta, thetaPriorDist, type = "l", col="green")
legend(0, 3.8, legend = c("prior", "posterior"), col = c("green", "red"), lty =
1:1, cex = 0.8)
```



b.)

Prior:

$$\theta \sim \text{Beta}(10, 2)$$

Posterior:

$$p(\theta | Y = 12) \propto p(\theta)p(Y = 12 | \theta) \propto \theta^9(1 - \theta)^{12}\theta^{12}(1 - \theta)^8 = \theta^{21}(1 - \theta)^9$$

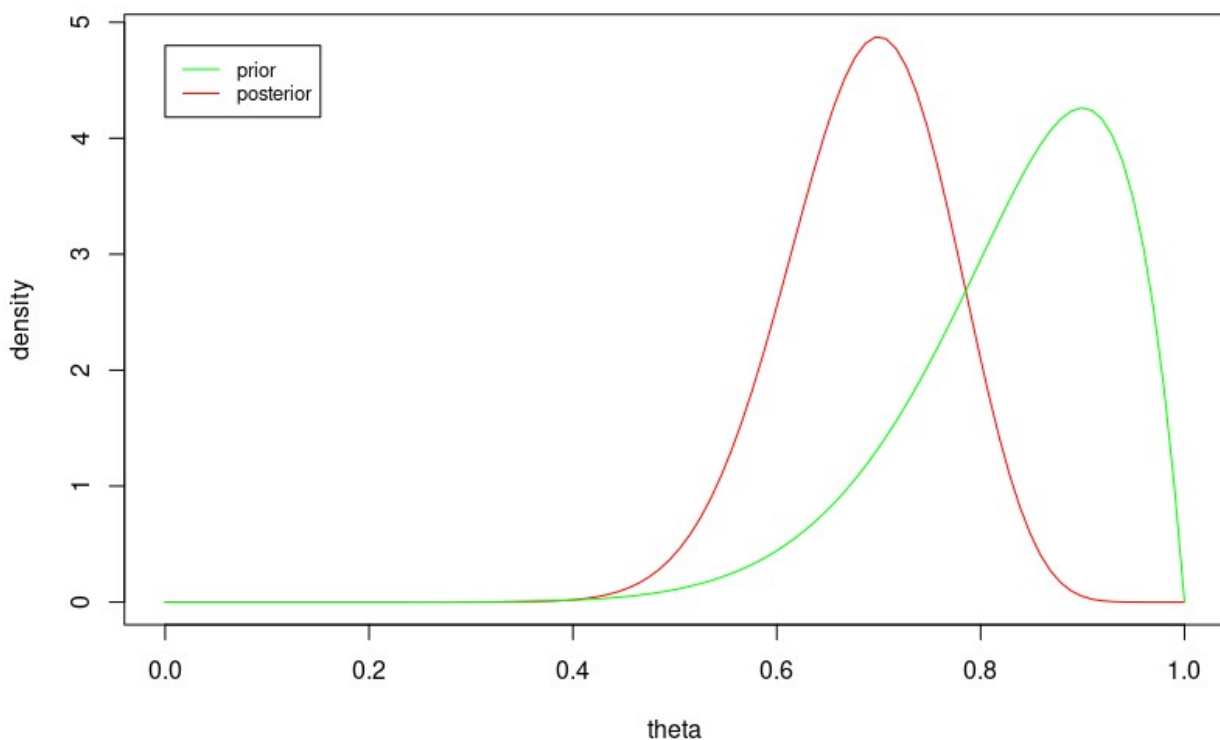
$$\therefore \theta | Y = 12 \sim \text{Beta}(22, 10)$$

$$\therefore E(\theta | Y = 12) = 0.6875, \text{Var}(\theta | Y = 12) = 0.0065$$

### R-Codes and plots:

```
data <- c(rep(1,12), rep(0,8))
thetaPostInfo <- bb.sum(data, 1, 10, 2)
thetaPostDist <- dbeta(theta, thetaPostInfo$alfa1, thetaPostInfo$beta1)
thetaPriorDist <- dbeta(theta, 10, 2)

plot(theta, thetaPostDist, type = "l", col="red", ylab="density")
lines(theta, thetaPriorDist, type = "l", col="green")
legend(0, 4.8, legend = c("prior", "posterior"), col = c("green", "red"), lty =
1:1, cex = 0.8)
```



c.)

In case a.) the non-informative prior distribution is used, so our posterior model is completely data driven. Our question in hand was formed as; Is an annual rate of degradation of gravestones more than 0.12mm, where  $\theta$  is the probability of this phenomenon. Based on the first model we could estimate  $\theta$  to fall in range [0.42, 0.76] with 90% probability. Intuitively this range seems a kind of “neutral” and so, in this specific graveyard, the gravestones doesn’t degrade significantly fast.

In case b.) the experimenters had some prior information about  $\theta$ , which can be modelled as Beta(10, 2) distribution. Based on the second plot, it can be seen that the prior distribution is strongly skewed to larger values of  $\theta$ , so compared to case a.) we might expect a bit more “extreme” posterior inference about  $\theta$ , especially because the sample size is small. Based on the posterior model we could estimate  $\theta$  to fall in range [0.55, 0.81] with 90% probability, so we get a bit more significant results compared to case a.). It might be a reasonable to say that in this particular cemetery, it is more likely that gravestones degrades more than 12mm per year in most cases. If we want to made more detailed conclusions, we might want to take a look to methods how experimenters constructed their prior believes in a form of Beta(10, 2), because in this case that distribution plays significant role also in posterior inference.

2.

a.)

Likelihood:

$$p(\mathbf{Y}_A | \lambda_A) = \prod_{i=1}^{10} p(y_{Ai} | \lambda_A) = \prod_{i=1}^{10} \lambda_A^{y_{Ai}} e^{-\lambda_A} \frac{1}{y_{Ai}!} = \lambda_A^{\sum y_{Ai}} e^{-10\lambda_A} \prod_{i=1}^{10} \frac{1}{y_{Ai}!}$$

Prior:

$$\lambda_A \sim \text{Gamma}(120, \frac{1}{10})$$

Posterior:

$$p(\lambda_A | \mathbf{Y}_A) \propto p(\lambda_A) p(\mathbf{Y}_A | \lambda_A) \propto \lambda_A^{120-1} e^{-10\lambda_A} \lambda_A^{\sum y_i} e^{-10\lambda_A} = \lambda_A^{120+\sum y_i-1} e^{-20\lambda_A}$$

$$\therefore \lambda_A | \mathbf{Y}_A \sim \text{Gamma}(120 + \sum_{i=1}^{10} y_{Ai}, \frac{1}{20}) \equiv \text{Gamma}(237, \frac{1}{20})$$

$$\therefore E(\lambda_A | \mathbf{Y}_A) = 11.85, \text{Var}(\lambda_A | \mathbf{Y}_A) = 0.5925$$

**R-codes:**

```
dataA <- c(12,9,12,14,13,13,15,8,15,6)
theta <- seq(7,17,length.out = 200)
thetaAPriorAlfa <- 120
thetaAPriorBeta <- 1/10 # 1/10
thetaAPriorDist <- dgamma(theta, thetaAPriorAlfa, scale=thetaAPriorBeta)
thetaAPostInfo <- pg.sum(dataA, thetaAPriorAlfa, 1/thetaAPriorBeta)
thetaAPostDist <- dgamma(theta, thetaAPostInfo$alfa1, rate=thetaAPostInfo$beta1)
thetaAPostMean <- thetaAPostInfo$Mean
thetaAPostVar <- thetaAPostInfo$Var
thetaAPostHDI <- pg.hdi(dataA, thetaAPriorAlfa, thetaAPriorBeta, 0.9)

cat("Theta A:",
    "\nPosterior mean:", thetaAPostMean,
    "\nPosterior variance:", thetaAPostVar,
    "\n90% HDI: [", thetaAPostHDI$Ala, ", ", thetaAPostHDI$Yla, "]\n")
```

Code output:

Theta A:

Posterior mean: 11.85

Posterior variance: 0.5925

90% HDI: [ 10.58073 , 13.10933 ]

b.)

Likelihood:

$$p(\mathbf{Y}_B | \lambda_B) = \lambda_B^{\sum y_{Bi}} e^{-13\lambda_B} \prod_{i=1}^{13} \frac{1}{y_{Bi}!}$$

Prior:

$$\lambda_B \sim \text{Gamma}(12, 1)$$

Posterior:

$$p(\lambda_B | \mathbf{Y}_B) \propto p(\lambda_B) p(\mathbf{Y}_B | \lambda_B) \propto \lambda_B^{12-1} e^{-\lambda_B} \lambda_B^{\sum y_{Bi}} e^{-13\lambda_B} = \lambda_B^{12+\sum y_{Bi}-1} e^{-14\lambda_B}$$

$$\therefore \lambda_B | \mathbf{Y}_B \sim \text{Gamma}(12 + \sum_{i=1}^{13} y_{Bi}, \frac{1}{14}) \equiv \text{Gamma}(125, \frac{1}{14})$$

$$\therefore E(\lambda_B | \mathbf{Y}_B) = 8.93, \text{Var}(\lambda_B | \mathbf{Y}_B) = 0.6378$$

**R-codes:**

```
dataB <- c(11,11,10,9,9,8,7,10,6,8,8,9,7)
theta <- seq(0,20,length.out = 200)
thetaBPriorAlfa <- 12
thetaBPriorBeta <- 1
thetaBPriorDist <- dgamma(theta, thetaBPriorAlfa, thetaBPriorBeta)
thetaBPostInfo <- pg.sum(dataB, thetaBPriorAlfa, thetaBPriorBeta)
thetaBPostDist <- dgamma(theta, thetaBPostInfo$alfa1, thetaBPostInfo$beta1)
#1/thetaBPostInfo$beta1
thetaBPostMean <- thetaBPostInfo$Mean
thetaBPostVar <- thetaBPostInfo$Var
thetaBPostHDI <- pg.hdi(dataA, thetaBPriorAlfa, thetaBPriorBeta, 0.9)

cat("Theta B:",
    "\nPosterior mean:", thetaBPostMean,
    "\nPosterior variance:", thetaBPostVar,
    "\n90% HDI: [", thetaBPostHDI$Alfa, ", ", thetaBPostHDI$Yla, "]\n")
```

Code output:

Theta B:

Posterior mean: 8.928571

Posterior variance: 0.6377551

90% HDI: [ 10.02459 , 13.41231 ]

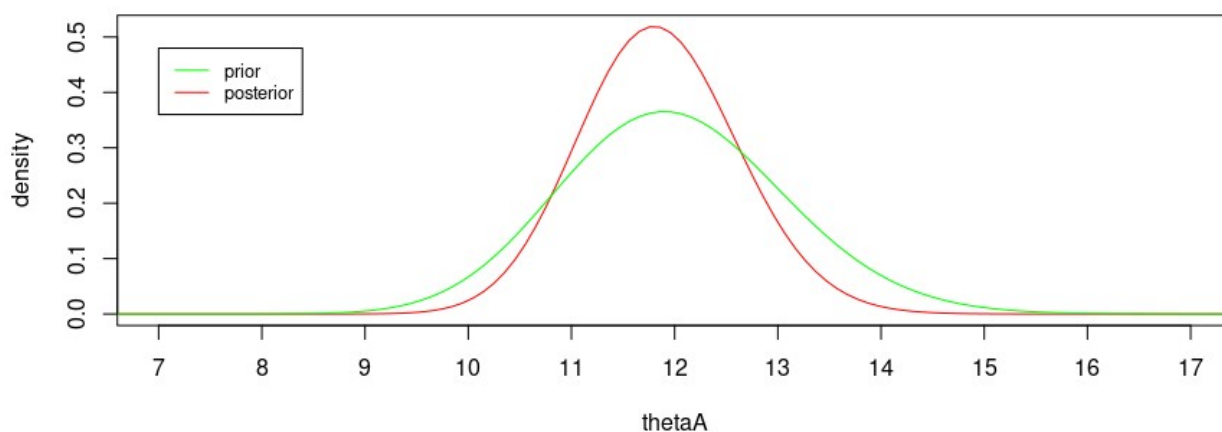
c.)

### R-codes and plots:

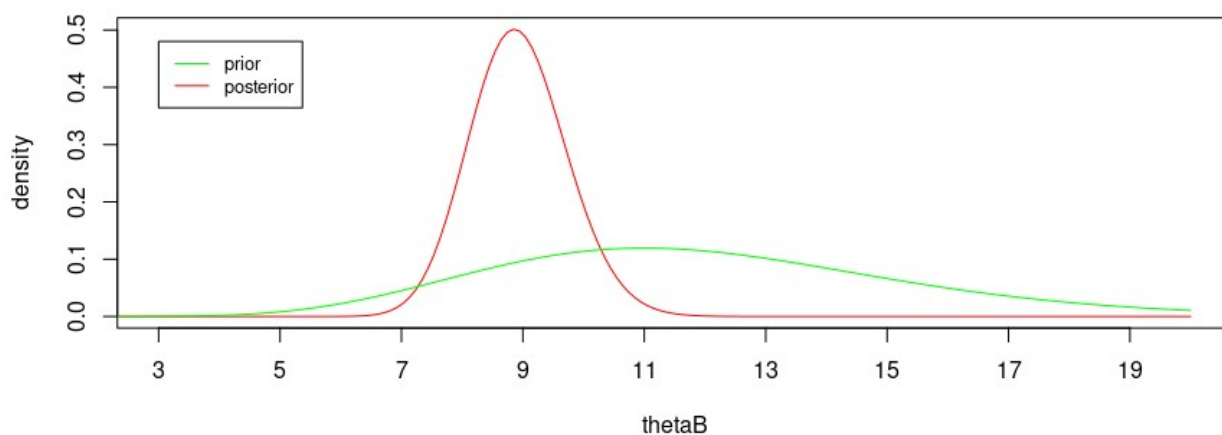
```
par(mfrow=c(2,1))
# Plotting thetaA
xaxis <- seq(7, 17, by = 1)
plot(theta, thetaAPostDist, type = "l", col = "red", ylab = "density", xlab =
"thetaA", xlim = c(7, 17), xaxt = 'n', main = "a.)")
axis(1, at = xaxis, las = 1)
lines(theta, thetaAPriorDist, type = "l", col = "green")
legend(7, 0.48, legend = c("prior", "posterior"), col = c("green", "red"), lty =
1:1, cex = 0.8)

# Plotting thetaB
xaxis <- seq(3, 20, by = 2)
plot(theta, thetaBPostDist, type = "l", col = "red", ylab = "density", xlab =
"thetaB", xlim = c(3, 20), xaxt = 'n', main = "b.)")
axis(1, at = xaxis, las=1)
lines(theta, thetaBPriorDist, type = "l", col = "green")
legend(3, 0.48, legend = c("prior", "posterior"), col = c("green", "red"), lty =
1:1, cex = 0.8)
```

a.)



b.)



d.)

**R-codes:**

```
dataB <- c(11,11,10,9,9,8,7,10,6,8,8,9,7)
#theta <- seq(0,20,length.out = 200)
postMeans <- data.frame(Prior_alpha=integer(0), Prior_beta=integer(0),
Post_alpha=integer(0), Post_beta=double(0), Post_mean=double(0))
for(n in 1:50){
  thetaBPriorAlfa <- 12*n
  thetaBPriorBeta <- n
  thetaBPostInfo <- pg.sum(dataB, thetaBPriorAlfa, 1/thetaBPriorBeta)
  postMeans[nrow(postMeans)+1,] <- c(thetaBPriorAlfa, thetaBPriorBeta,
                                     thetaBPostInfo$alfa1,
                                     thetaBPostInfo$beta1,
                                     thetaBPostInfo$Mean)
}
postMeans
```

Code output:

	Prior_alpha	Prior_beta	Post_alpha	Post_beta	Post_mean
1	12	1	125	14.00000	8.928571
2	24	2	137	13.50000	10.148148
3	36	3	149	13.33333	11.175000
4	48	4	161	13.25000	12.150943
5	60	5	173	13.20000	13.106061
6	72	6	185	13.16667	14.050633
7	84	7	197	13.14286	14.989130
8	96	8	209	13.12500	15.923810
9	108	9	221	13.11111	16.855932
10	120	10	233	13.10000	17.786260
11	132	11	245	13.09091	18.715278
12	144	12	257	13.08333	19.643312
13	156	13	269	13.07692	20.570588
14	168	14	281	13.07143	21.497268
15	180	15	293	13.06667	22.423469
16	192	16	305	13.06250	23.349282
17	204	17	317	13.05882	24.274775
18	216	18	329	13.05556	25.200000
19	228	19	341	13.05263	26.125000
20	240	20	353	13.05000	27.049808
21	252	21	365	13.04762	27.974453
22	264	22	377	13.04545	28.898955
23	276	23	389	13.04348	29.823333
24	288	24	401	13.04167	30.747604
25	300	25	413	13.04000	31.671779
26	312	26	425	13.03846	32.595870
27	324	27	437	13.03704	33.519886
28	336	28	449	13.03571	34.443836
29	348	29	461	13.03448	35.367725
30	360	30	473	13.03333	36.291560
31	372	31	485	13.03226	37.215347
32	384	32	497	13.03125	38.139089
33	396	33	509	13.03030	39.062791
34	408	34	521	13.02941	39.986456
35	420	35	533	13.02857	40.910088
36	432	36	545	13.02778	41.833689
37	444	37	557	13.02703	42.757261

38	456	38	569	13.02632	43.680808
39	468	39	581	13.02564	44.604331
40	480	40	593	13.02500	45.527831
41	492	41	605	13.02439	46.451311
42	504	42	617	13.02381	47.374771
43	516	43	629	13.02326	48.298214
44	528	44	641	13.02273	49.221640
45	540	45	653	13.02222	50.145051
46	552	46	665	13.02174	51.068447
47	564	47	677	13.02128	51.991830
48	576	48	689	13.02083	52.915200
49	588	49	701	13.02041	53.838558
50	600	50	713	13.02000	54.761905

It can be seen that the posterior expectation of  $\theta_B$  is closest to that of to be  $\theta_A$ , when the posterior distribution is formed as (output line 4)

$$\theta_B | \mathbf{Y}_B \sim \text{Gamma}(161, \frac{1}{13.25})$$

and prior distribution for  $\theta_B$  is formed as

$$\theta_B \sim \text{Gamma}(48, 4)$$

To get the prior distribution for  $\theta_B$ , we have to examine and reorganize the new posterior pdf:

$$p(\theta_B | \mathbf{Y}_B) \propto \theta_B^{48 + \sum y_{Bi} - 1} e^{-\theta_B(\frac{1}{4} + 13)} \propto \theta_B^{48-1} e^{-\frac{1}{4}\theta_B} \theta_B^{\sum y_{Bi}} e^{-13\theta_B}$$

Thus our updated prior believes about  $\theta_B$  can be modelled by distribution:

$$\theta_B \sim \text{Gamma}(48, 4)$$



3.)

a.)

### R-codes:

```
fn <- "/home/tuomas/R/Projects/Bayesian analysis/Take_home_task/task3.jag"
cat(
  "model
  {
    # Likelihood:
    for(i in 1:N)
    {
      x[i] ~ dnegbin(p[i],1)
      # Prior for p's
      p[i] ~ dbeta(a, b)
    }
    # Hyperpriors for a and b
    a ~ dgamma(1, 1)
    b ~ dgamma(1, 1)
    #b ~ dgamma(1, 0.5)

    theta <- a/(a+b)
  }
  ",
  file = fn )
```

```
task3.data <- list(x=c(4,0,1,7,3,2,8,0), N=8)
task3.init <- list(p=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5), a=11, b=1)
task3.model <- jags.model(file = fn, data = task3.data, inits = task3.init,
n.chains = 3, n.adapt = 2000)
```

```
task3.par <- c("p[1]", "p[2]", "p[3]", "p[4]", "p[5]", "p[6]", "p[7]", "p[8]", "theta")
task3.result <- coda.samples(model = task3.model, variable.names = task3.par,
n.iter = 10000, thin = 10)
summary(task3.result)
```

#### Output:

1. Empirical mean and standard deviation for each variable,  
plus standard error of the mean:

	Mean	SD	Naive SE	Time-series SE
p[1]	0.2661	0.1547	0.002824	0.002825
p[2]	0.5494	0.2440	0.004455	0.004362
p[3]	0.4230	0.2143	0.003912	0.003836
p[4]	0.1928	0.1182	0.002157	0.002105
p[5]	0.3032	0.1703	0.003109	0.003049
p[6]	0.3581	0.1903	0.003475	0.003490
p[7]	0.1762	0.1105	0.002018	0.002018
p[8]	0.5431	0.2475	0.004519	0.004518
theta	0.3824	0.1087	0.001985	0.001985

2. Quantiles for each variable:

	2.5%	25%	50%	75%	97.5%
p[1]	0.03300	0.14723	0.2450	0.3633	0.6256
p[2]	0.10663	0.36323	0.5497	0.7388	0.9774
p[3]	0.06973	0.25587	0.4030	0.5709	0.8590
p[4]	0.02422	0.10579	0.1730	0.2620	0.4709

```
p[5] 0.04192 0.17055 0.2817 0.4185 0.6797
p[6] 0.04959 0.20897 0.3376 0.4905 0.7604
p[7] 0.02223 0.09242 0.1559 0.2407 0.4362
p[8] 0.10373 0.34988 0.5397 0.7438 0.9762
theta 0.19396 0.30421 0.3742 0.4562 0.6130
```

Mean and standard deviation for  $p_i$ 's are listed in output table 1  
 95% credible interval for  $p_i$ 's can be found in output table 2 so that  
 CI = [2.5% quantile, 97.5% quantile]

b.)

Code output when  $b \sim \text{Gamma}(1,2)$ :

1. Empirical mean and standard deviation for each variable,  
 plus standard error of the mean:

	Mean	SD	Naive SE	Time-series SE
p[1]	0.2524	0.1430	0.002610	0.002597
p[2]	0.4707	0.2346	0.004283	0.004372
p[3]	0.3925	0.1968	0.003593	0.003591
p[4]	0.1917	0.1126	0.002055	0.001967
p[5]	0.2910	0.1604	0.002929	0.003109
p[6]	0.3294	0.1772	0.003234	0.003215
p[7]	0.1771	0.1088	0.001987	0.001928
p[8]	0.4752	0.2355	0.004300	0.004138

2. Quantiles for each variable:

	2.5%	25%	50%	75%	97.5%
p[1]	0.03867	0.14244	0.2322	0.3433	0.5793
p[2]	0.08765	0.28451	0.4492	0.6400	0.9487
p[3]	0.07190	0.24255	0.3724	0.5267	0.8186
p[4]	0.02832	0.10608	0.1729	0.2575	0.4539
p[5]	0.04844	0.16584	0.2696	0.3964	0.6437
p[6]	0.05440	0.19485	0.3055	0.4536	0.7109
p[7]	0.02465	0.09431	0.1566	0.2425	0.4357
p[8]	0.08782	0.29320	0.4516	0.6496	0.9421

c.)

Posterior inference for  $\theta = \frac{\alpha}{\alpha+\beta}$ :

Empirical mean and standard deviation:

	Mean	SD	Naive SE	Time-series SE
theta	0.3819	0.1086	0.001984	0.001997

Quantiles:

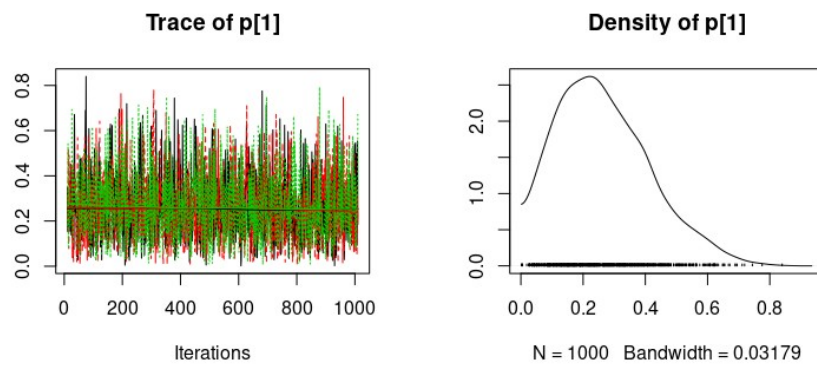
	2.5%	25%	50%	75%	97.5%
theta	0.18885	0.30576	0.3740	0.4525	0.6144

d.)

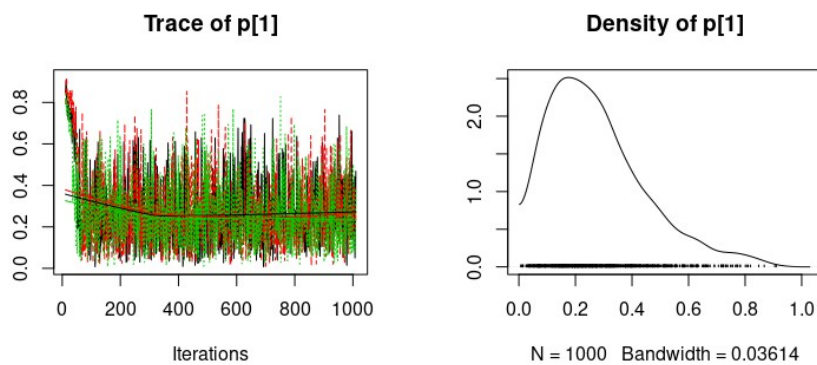
The choose of initial values for markov chain has effect only on so called burn-in-phase which is an early stage of simulated mcmc process. If the initial values are poorly chosen, it will take more time to markov chain to converge towards posterior distribution, and thus decreases the performance, which for example, can be seen from the trace-plot of mcmc process. There are quite a few of ways to select initial values for parameters, but some preferred methods are to use MLE-estimate of a parameter, or alternatively, when using multiple chains it is often preferred to start from dispersed points in parameter space. Below are listed traceplots using various initalization values. Number of

iterations were decreased to 1000, so that the burn-in phase could be seen more easily and n.adapt is set to 10 so we can inspect the early stage of markov chain more clearly.

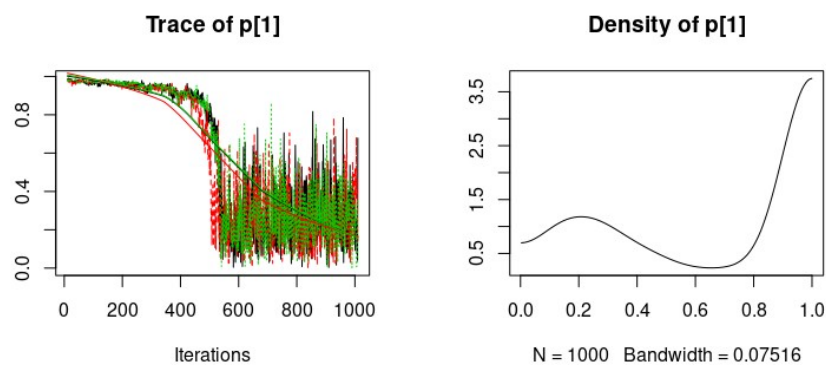
Case 1:  $a=11$ ,  $b=1$ ,  $n.adapt=10$ ,  $n.iter=1000$ :



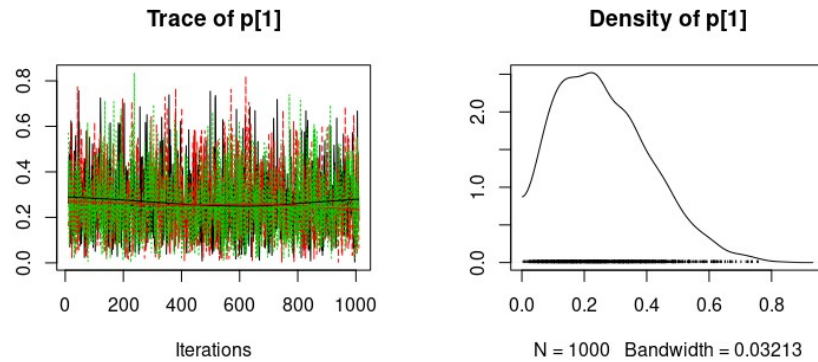
Case 2:  $a=111$ ,  $b=1$ ,  $n.adapt=10$ ,  $n.iter=1000$ :



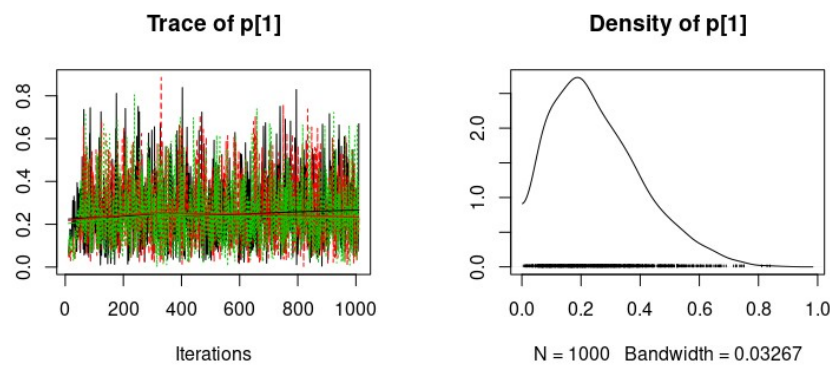
Case 3:  $a=1111$ ,  $b=1$ ,  $n.adapt=10$ ,  $n.iter=1000$ :



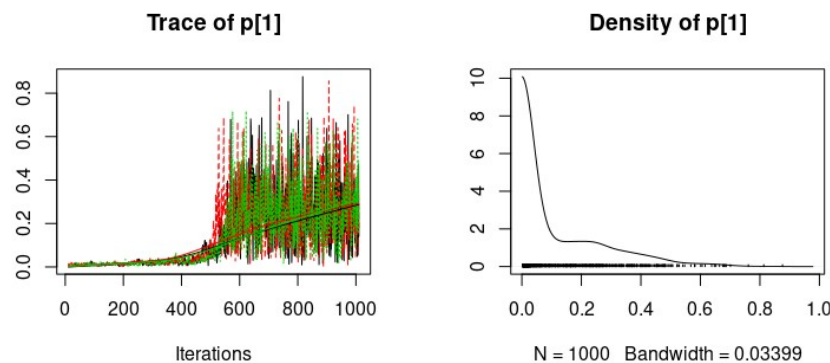
Case 4:  $a=1$ ,  $b=11$ ,  $n.adapt=10$ ,  $n.iter=1000$ :



Case 5:  $a=1$ ,  $b=111$ ,  $n.adapt=10$ ,  $n.iter=1000$ :



Case 6:  $a=1$ ,  $b=1111$ ,  $n.adapt=10$ ,  $n.iter=1000$ :



As can be seen, the effect of extreme parameter values dominates the early stage of mcmc process. However, in these examples only the first 10 samples are discarded from markov chain and a result from this will produce highly biased posterior information. If  $n.adapt$  is set as a larger value, the poorly selected parameter values doesn't have that much effect to the posterior inference.

4.

**R-codes:**

```
fn <- "/home/tuomas/R/Projects/Bayesian analysis/Take_home_task/task4.jag"
cat("
  model{
    # Likelihood:
    for(i in 1:N){
      mu[i] <- b0 + b1*weight[i] + b2*sixcyl[i] + b3*eightcyl[i] + e
      mpg[i] ~ dnorm(mu[i], prec)
    }
    # Priors:
    b0 ~ dnorm(0, 0.0001) # Variance = 1000
    b1 ~ dnorm(0, 0.0001)
    b2 ~ dnorm(0, 0.0001)
    b3 ~ dnorm(0, 0.0001)
    e ~ dnorm(0, prec)
    prec ~ dgamma(0.01, 0.01) #?
    sigma2 <- 1/sqrt(prec)

    # Predictive distribution:
    mu.new <- b0 + b1*weight.new + e
    mpg.new ~ dnorm(mu.new, prec)

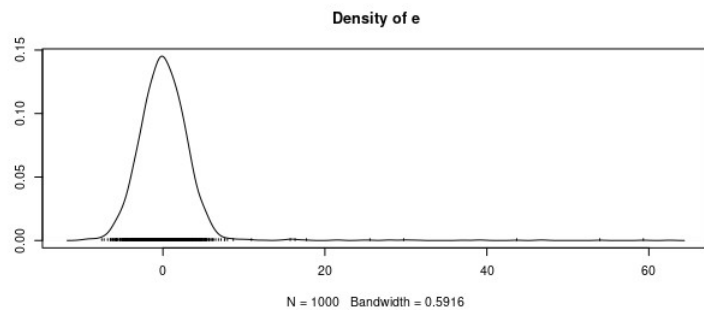
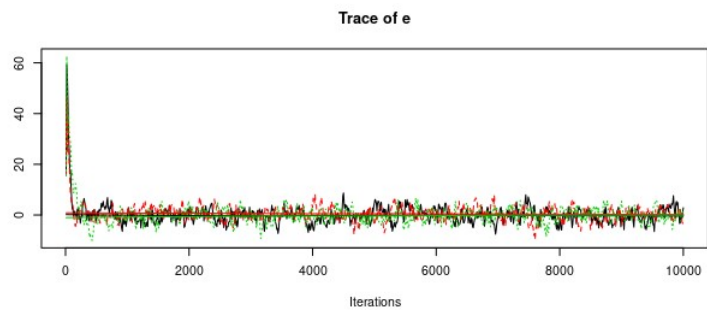
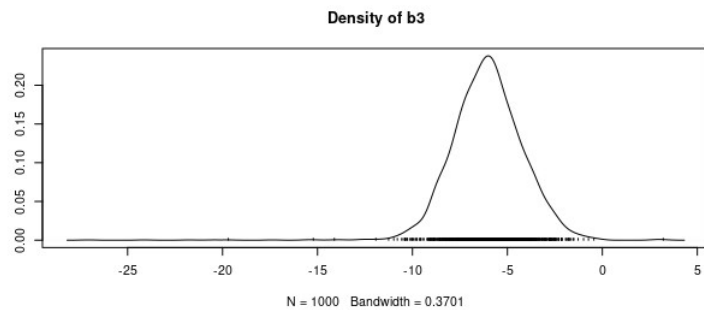
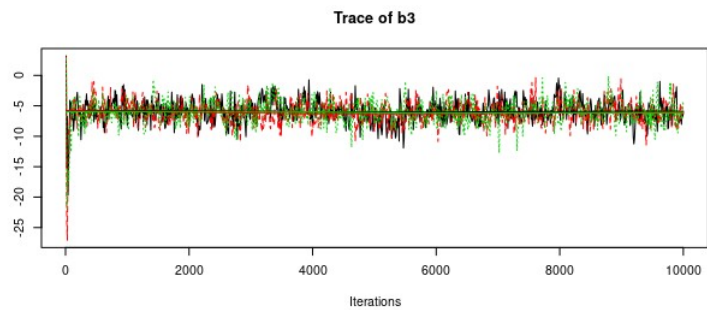
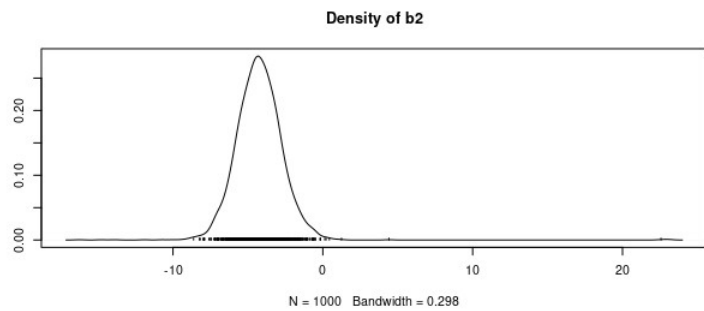
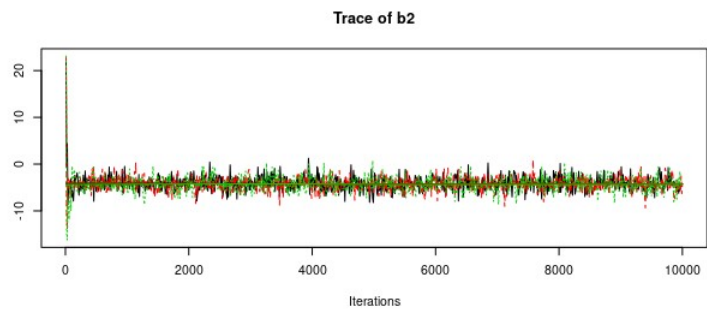
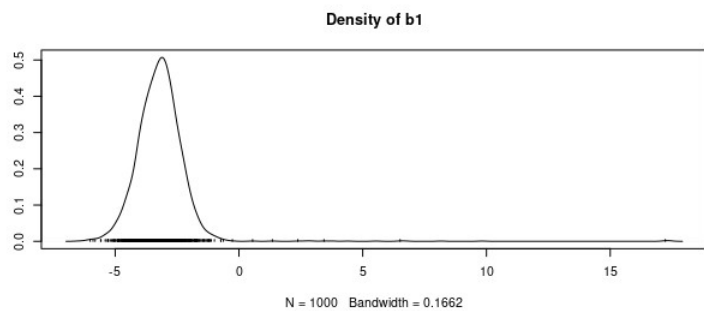
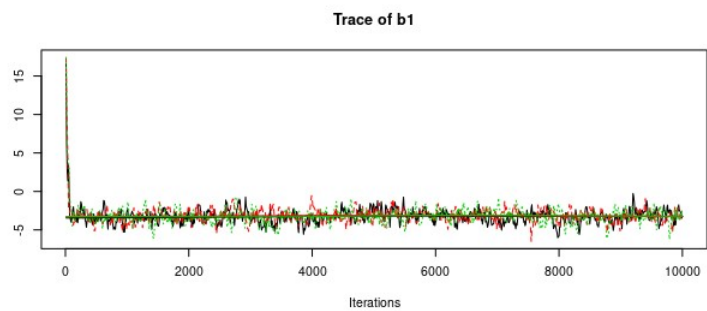
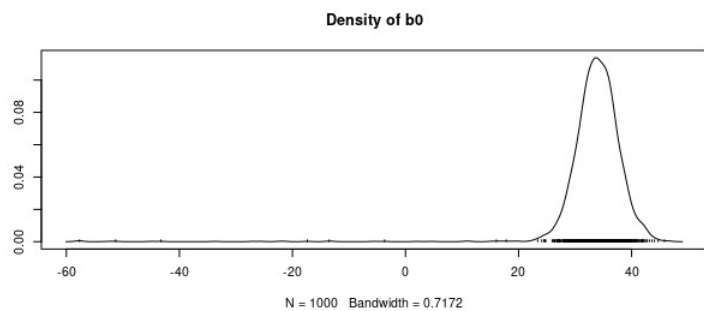
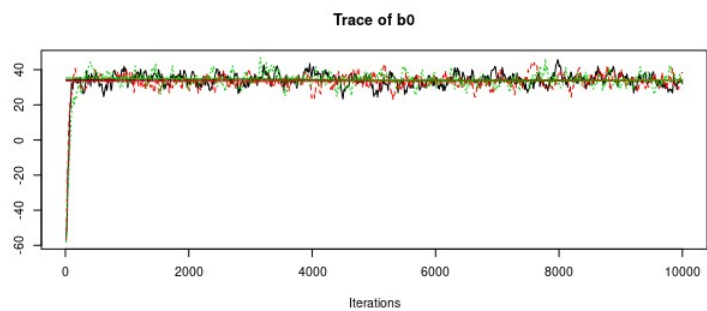
  }",
  file = fn )

task4.data <- list(mpg=c(21.0,21.0,22.8,21.4,18.7,18.1,14.3,24.4,22.8,19.2,17.8,
                        16.4,17.3,15.2,10.4,10.4,14.7,32.4,30.4,33.9,21.5,15.5,
                        15.2,13.3,19.2,27.3,26.0,30.4,15.8,19.7,15.0, 21.4),
                  weight=c(2.620, 2.875, 2.320, 3.215, 3.440, 3.460, 3.570,
                           3.190, 3.150, 3.440, 3.440, 4.070, 3.730, 3.780,
                           5.250, 5.424, 5.345, 2.200, 1.615, 1.835, 2.465,
                           3.520, 3.435, 3.840, 3.845, 1.935, 2.140, 1.513,
                           3.170, 2.770, 3.570, 2.780),
                  sixcyl=c(1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
                           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0),
                  eightcyl=c(0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1,
                              0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0),
                  weight.new = 3.5,
                  mpg.new = NA,
                  N=32
                  )

task4.init <- list(b0=20,b1=20,b2=20,b3=20,e=0.5,prec=1)
task4.model <- jags.model(fn, data = task4.data, inits = task4.init,
                          n.chains = 3, n.adapt = 2000)

task4.pars <- c("b0", "b1", "b2", "b3", "e")
task4.postInfo <- coda.samples(model = task4.model, variable.names = task4.pars,
                              n.iter = 10000, thin = 10)

summary(task4.postInfo)
plot(task4.postInfo)
```



Console output:

Iterations = 10:10000  
Thinning interval = 10  
Number of chains = 3  
Sample size per chain = 1000

1. Empirical mean and standard deviation for each variable,  
plus standard error of the mean:

	Mean	SD	Naive SE	Time-series SE
b0	33.6046	6.332	0.11560	0.33247
b1	-3.1780	1.168	0.02133	0.04888
b2	-4.2375	1.723	0.03145	0.03366
b3	-6.0599	1.939	0.03540	0.06144
e	0.2721	4.143	0.07563	0.22163

2. Quantiles for each variable:

	2.5%	25%	50%	75%	97.5%
b0	26.418	31.766	33.96030	36.262	41.099
b1	-4.902	-3.749	-3.20044	-2.707	-1.583
b2	-7.093	-5.194	-4.28532	-3.326	-1.235
b3	-9.527	-7.208	-6.05986	-4.887	-2.471
e	-5.305	-1.806	0.01488	1.903	5.709

b.)

95% equal tail probability intervals:

$\beta_0$  : [27.004, 40.890]  
 $\beta_1$  : [-4.858, -1.584]  
 $\beta_2$  : [-7.170, -1.315]  
 $\beta_3$  : [-9.515, -2.502]  
*error* : [-5.464, 5.925]

c.)

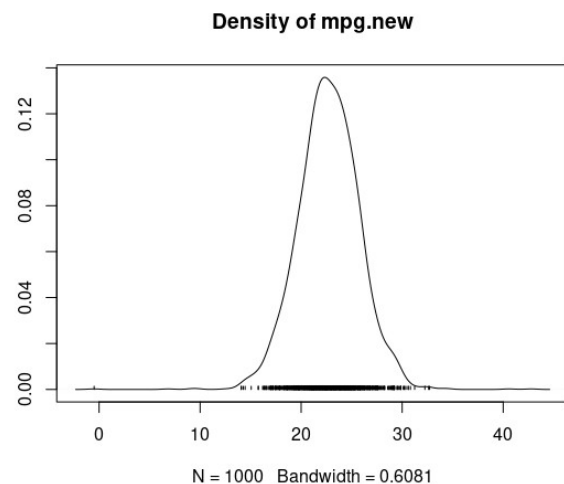
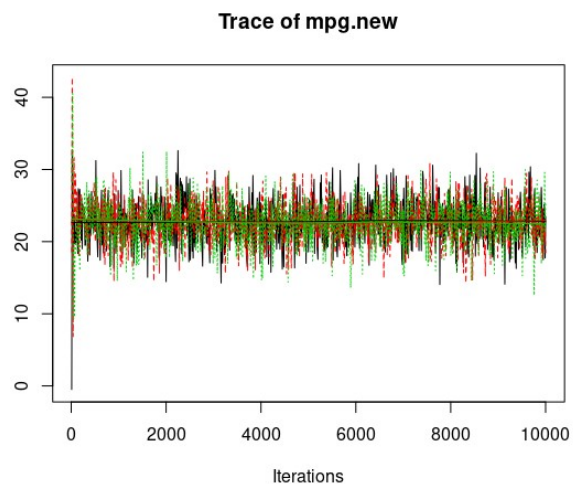
There is not significant difference in mpg when comparing six and eight cylinder cars. When viewing the probability intervals in task b.) it can be seen that the intervals for  $\beta_2$  and  $\beta_3$  extensively overlaps each other. Taking into account also the dummy binary variables, the overall difference is not notable.

d.)

### R-codes and plots:

```
task4d.pars <- c("mpg.new")
task4d.postInfo <- coda.samples(model = task4d.model,
                                variable.names = task4d.pars,
                                n.iter = 10000, thin = 10)

summary(task4d.postInfo)
plot(task4d.postInfo)
```



Console output:

1. Empirical mean and standard deviation for each variable (mpg.new), plus standard error of the mean:

Mean	SD	Naive SE	Time-series SE
22.89788	3.09286	0.05647	0.06324

2. Quantiles for each variable (mpg.new):

2.5%	25%	50%	75%	97.5%
17.00	20.91	22.85	24.83	28.93



5.

a.)

### R-codes and plots:

```
cat("
  model{
    # Likelihood:
    for(i in 1:N){
      y1[i] ~ dnorm(mu, sigma1)
      y2[i] ~ dnorm(mu, sigma2)

    }
    # Priors:
    mu ~ dnorm(4.7, 1/v)
    sigma1 ~ dgamma(4, 1/5)
    sigma2 ~ dgamma(4, 1/5)
    tau1 <- 1/sigma1
    tau2 <- 1/sigma2
    phi <- tau2/tau1
  }",
  file = fn)

# a)
task5.data <- list(y1 = c(4.3, 4.3, 2.7, 3.6, 3.5, 4.5),
                  y2 = c(3.9, 4.0, 4.5, 2.9, 5.2, 4.8),
                  v = 0.2,
                  N = 6)
task5.init <- list(mu=4.7, sigma1=1, sigma2=1)
task5.model <- jags.model(file = fn, data = task5.data, inits = task5.init,
n.chains = 3, n.adapt = 2000)

task5.pars <- c("mu", "phi", "tau1", "tau2")
task5.postInfo <- coda.samples(model = task5.model, variable.names = task5.pars,
                              n.iter = 100000, thin = 10)

summary(task5.postInfo)
plot(task5.postInfo)
```

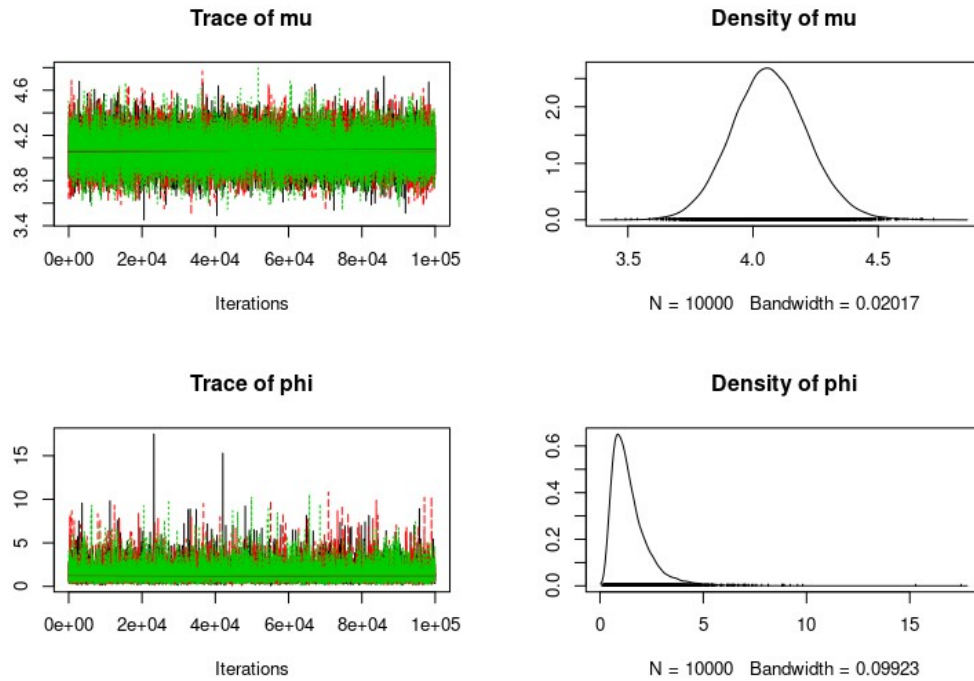
Code output:

1. Empirical mean and standard deviation for each variable,  
plus standard error of the mean:

	Mean	SD	Naive SE	Time-series SE
mu	4.0682	0.1520	0.0008777	0.0008777
phi	1.4554	0.9384	0.0054178	0.0054354
tau1	0.2716	0.1312	0.0007572	0.0007676
tau2	0.3294	0.1493	0.0008619	0.0008619

2. Quantiles for each variable:

	2.5%	25%	50%	75%	97.5%
mu	3.7774	3.9651	4.0661	4.1685	4.3731
phi	0.3838	0.8260	1.2237	1.8127	3.8848
tau1	0.1170	0.1850	0.2414	0.3224	0.6083
tau2	0.1491	0.2305	0.2951	0.3896	0.7139



Inference about  $\mu < 4.7$  :

Based on 95% CI [3.7785, 4.377] of  $\mu$ , it can be inferred that it is quite likely that  $\mu < 4.7$

Inference about  $\phi = \frac{\tau_2}{\tau_1}$  :

Based on 95% credible interval and posterior mean of  $\phi$  it may be estimated that  $\tau_2$  is roughly equal with  $\tau_1$ . However, it must also be noted that the sample size was rather small and due to this, the prior distributions of  $\tau$ 's will be in dominant role which in this case, are equal. In this sense it may be reasonable to say, that  $\tau_2 > \tau_1$  because, although the sample is small, there is still slight difference between  $\tau$ 's posterior information, even though their prior distributions are equal.

b.)

When  $v$  grows indefinitely, the prior distribution of  $\mu$  becomes more and more noninformative and thus the posterior distribution of  $\mu$  becomes more and more data driven, a.k.a. the posterior mean of  $\mu$  converges towards sample mean. This can be noted when sampling from posterior distribution at the same time when increasing values for  $v$ . Below are some numbers regarding to this method:

	$v$	$\mu_{\text{mean}}$
1	1.500000e-01	4.092256
2	5.000000e-01	4.023562
3	1.000000e+00	4.006566
4	5.000000e+00	3.991787
5	1.000000e+01	3.991407
6	5.994843e+01	3.988174
7	3.593814e+02	3.988363
8	2.154435e+03	3.988377
9	1.291550e+04	3.988715
10	7.742637e+04	3.987182
11	4.641589e+05	3.987386
12	2.782559e+06	3.988744
13	1.668101e+07	3.988994
14	1.000000e+08	3.989758

c.)

Based on examination, it can be noted that the new process produces less wastage, compared to the standard process. Average wastage from standard process was measured to be 4.7% but by using updated process, the average wastage percentage was measured to be 4.1%. To made a more specific summary about updated process, it was measured that by 95% probability the wastage percentage falls between 3.78% and 4.38%.

When comparing the variances, it may be reasonable to assume that the variance regarding to operator 2 might be slightly larger than variance regarded to operator 1. However, the sample of given measurement data was rather small. To obtain more secure inferences about variances, more measurement data would be needed.